



# Pintos Project 1: Threads

Hoonsung Chwa

Minsup Lee

Juyoung Lee

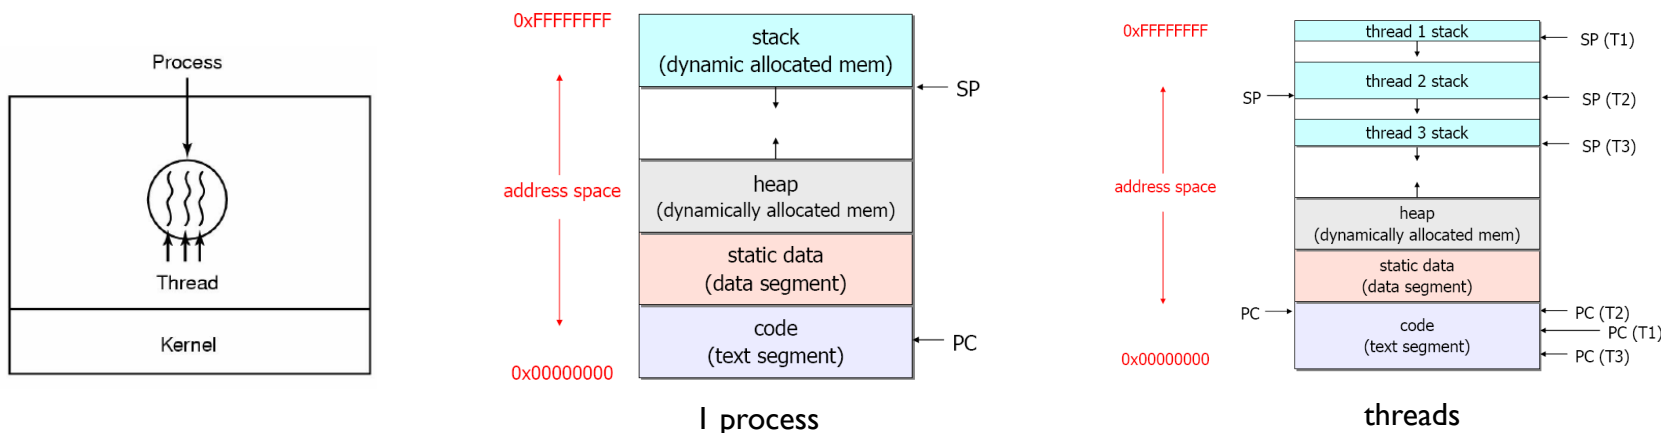
Division of Computer Science, KAIST

**KAIST**

# What is a Thread?

## ■ Thread

- A flow of control in a process
- A sequence of instructions being executed in a program.
- Threads share the process instructions and most of its data.
  - A change in shared data by one thread can be seen by the other threads in the process



From Prof. Jin-Soo Kim's slides

# To do

- **Alarm Clock**
- **Priority Scheduler**
- **Advanced Scheduler (Optional)**
- **Design Document**

# Alarm Clock (1)



## ▪ To suspend threads for some timer ticks

- Threads call `timer_sleep(TICKS)` in `device/timer.c`
- Ex) `thread_create(thread_id, NULL, sleeper, ticks)`

```
/* Sleeper thread. */
static void
sleeper (void *t_)
{
    struct sleep_thread *t = t_;
    struct sleep_test *test = t->test;
    int i;

    for (i = 1; i <= test->iterations; i++)
    {
        int64_t sleep_until = test->start + i * t->duration;
        timer_sleep (sleep_until - timer_ticks ());

        lock_acquire (&test->output_lock);
        *test->output_pos++ = t->id;
        lock_release (&test->output_lock);
    }
}
```

# Alarm Clock (2)



## ■ Requirement

- Must not be "busy waiting"

## ■ Timer\_sleep() in device/timer.c

- Now, implemented by 'busy waiting'

```
/* Suspends execution for approximately TICKS timer ticks. */  
void  
timer_sleep (int64_t ticks)  
{  
    int64_t start = timer_ticks ();  
    while (timer_elapsed (start) < ticks)  
        thread_yield ();  
}
```

- Reimplement timer.c to avoid 'busy waiting'
  - Make a 'thread\_list' for waiting threads.
  - When timer interrupt occurs, wake the threads up
    - » Timer\_interrupt() in device/timer.c
    - » Ticks++

# Priority Scheduling (1)

## ■ Now, Round-Robin scheduling

- ready\_list in thread.c
- Schedule(), next\_thread\_to\_run() in thread.c

```
static void
schedule (void)
{
    struct thread *cur = running_thread ();
    struct thread *next = next_thread_to_run ();
    struct thread *prev = NULL;
    :
    if (cur != next)
        prev = switch_threads (cur, next);
    schedule_tail (prev);
}
```

```
static struct thread *
next_thread_to_run (void)
{
    if (list_empty (&ready_list))
        return idle_thread;
    else
        return list_entry (list_pop_front (&ready_list),
                           struct thread, elem);
}
```

## ■ Implement priority scheduling

- Thread which has the highest priority in the ready\_list gets CPU.

# Priority Scheduling (2)



## ■ Requirements

- If a thread is added to the ready list which has a higher priority than the currently running thread, **immediately** yield the processor to the new thread.
- Also, when threads are waiting for a lock, semaphore, or condition variable, the highest priority waiting thread should be woken up first.
- Solve the '**Priority Inversion**' problem of 'Lock'.
  - Implement priority inheritance (effective priority).

## ■ References

- Thread\_create(), next\_thread\_to\_run(), ... in thread.c
- Sema\_up(), lock\_acquire(), cond\_signal(), ... in synch.c



# Priority Scheduling (3)

## ■ Priority inversion problem


- A situation where a higher-priority job is unable to run because a lower-priority job is holding a resource it needs, such as a lock.

## ■ Example

- There are three threads, H, M and L
  - Priority:  $H > M > L$
- H and L need same resource (e.g. lock) and L holds the resource
  - Expected execution order is L, H, M
- In real, H is blocked and M is firstly executed. Then L and H.
  - $P(M) > P(L)$
- Consequently, the execution order is M, L, H.



# Priority Scheduling (4)



- **Input order : L, H, M**
- **Execution order : M, L, H**

From 2007 CS330's slides

# Priority Scheduling (5)




## ■ Priority inversion problem

- A situation where a higher-priority job is unable to run because a lower-priority job is holding a resource it needs, such as a lock.

## ■ Solution: Priority Inheritance


- The higher-priority job **donate** its priority to the lower-priority job holding the resource it requires (**effective priority**).
- The beneficiary of the inheritance will now have a higher scheduling priority.
  - Get scheduled to run sooner.
- It can finish its work and release the resource, at which point to the original priority is returned to the job.

# Priority Scheduling (6)




# Priority Scheduling (7)

- **Nested Donation**



- **Multiple Donation**



# Semaphore



## ■ synch.c and synch.h

```
/* A counting semaphore. */
struct semaphore
{
    unsigned value;           /* Current value. */
    struct list waiters;      /* List of waiting threads. */
};

void sema_init (struct semaphore *, unsigned value);
void sema_down (struct semaphore *);
bool sema_try_down (struct semaphore *);
void sema_up (struct semaphore *);
void sema_self_test (void);
```

## ■ Usage

- Sema\_down();
- critical section
- Sema\_up();

# Lock

## ■ synch.c and synch.h

```
/* Lock. */
struct lock
{
    struct thread *holder; /* Thread holding lock (for debugging). */
    struct semaphore semaphore; /* Binary semaphore controlling access. */
};

void lock_init (struct lock *);
void lock_acquire (struct lock *);
bool lock_try_acquire (struct lock *);
void lock_release (struct lock *);
bool lock_held_by_current_thread (const struct lock *);

static void
a_thread_func (void *lock_)
{
    struct lock *lock = lock_;

    lock_acquire (lock);
    msg ("Thread a acquired lock a.");
    lock_release (lock);
    msg ("Thread a finished.");
}
```

→ critical section

# Condition variable



## ■ **synch.c and synch.h**

```
/* Condition variable. */
struct condition
{
    struct list waiters;      /* List of waiting threads. */
};

void cond_init (struct condition *);
void cond_wait (struct condition *, struct lock *);
void cond_signal (struct condition *, struct lock *);
void cond_broadcast (struct condition *, struct lock *);
```

## ■ **Usage**

- Thread1
  - Cond\_wait(&cv, &lock\_cv);
- Thread2
  - Cond\_signal(&cv, &lock\_cv); or cond\_broadcast(&cv, &lock\_cv);
- A lock is required to prevent race condition.



# Advanced Scheduler



- It is optional.
- If you implement *advanced scheduler(4.4BSD scheduler)*, you can get extra score.
- Read the below link carefully
  - [http://www.scs.stanford.edu/07au-cs140/pintos/pintos\\_7.html#SEC128](http://www.scs.stanford.edu/07au-cs140/pintos/pintos_7.html#SEC128)

# Submission



- **Due**

- 10/4 (Sun) 24:00

- **Help desk**

- 9/24 (Thur) 19:00~22:00 at Eve room(1<sup>st</sup> floor)

- **E-mail to 'cs330\_submit AT calab.kaist.ac.kr'**

- Title: [cs330] project 1
- Contents
  - Source code (archive of 'pintos/src' directory)
  - Design Documentation
    - » [http://cps.kaist.ac.kr/courses/2009\\_fall\\_cs330/project/project1.tmpl](http://cps.kaist.ac.kr/courses/2009_fall_cs330/project/project1.tmpl)

- **Cheating will not be forgiven.**

# Tips

- **Strongly RECOMMEND to use 'ctags'**
- **Don't make list/hash structure.**
  - Use list.h list.c, hash.h hash.c at lib/kernel
- **Check progress yourself by using 'make check'**
  - alarm-single, alarm-multiple, alarm-simultaneous, alarm-priority, alarm-zero, alarm-negative,
  - priority-change, priority-donate-one, priority-donate-multiple, priority-donate-multiple2, priority-donate-nest, priority-donate-sema, priority-donate-lower, priority-fifo, priority-preempt, priority-sema, priority-condvar, priority-donate-chain
  - mlfqs-load-1, mlfqs-load-60, mlfqs-load-avg, mlfqs-recent-1, mlfqs-fair-2, mlfqs-fair20, mlfqs-nice-2, mlfqs-nice10, mlfqs-block
- **TAs**
  - Noah course board
  - 좌훈승, chwahs\_at\_gmail.com
  - 이민섭, ssangbujja\_at\_cps.kaist.ac.kr
  - 이주영, lly8904\_at\_kaist.ac.kr